

Multiple linear regression

Introduction

In multiple linear regression, we have a dependent variable Y which we would like to explain in terms of the multiple independent variables X_1, X_2, \dots, X_q simultaneously.




Figure 1: Figure 1: A multiple regression model (with two explanatory variables) represented as a hyperplane in three-dimensional space. The dots surrounding the hyperplane represent the observed data.

The general notation of the model is:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_q X_q + \epsilon$$

where

Y is the value of the response variable;

X_j is the value of the j explanatory variable;

β_0 is the intercept, and β_j is the partial slope of the j explanatory variable;

ϵ is that part of Y which is not explained by the regression (i.e. the error).

The parameters, $\beta_0, \beta_1, \beta_2, \dots, \beta_q$, must be estimated from the data. The parameter estimates, $\hat{\beta}_0, \hat{\beta}_1, \hat{\beta}_2, \dots, \hat{\beta}_q$, are estimated by R using the method of least squares.

The formula of the fitted model is: $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_q X_q$

Hypotheses concerning the parameters can be tested. For example, for a particular $j = 1, \dots, q$:

0: $\beta_j = 0$ (X_j does not “contribute” to the explanation of Y)

1: $\beta_j \neq 0$ (X_j does “contribute” to the explanation of Y)

Load required packages

```
library(RcmdrMisc)
```

Loading required package: car

Loading required package: carData

Loading required package: sandwich

```
library(ggplot2)
library(car)
library(MASS)
```

Importing data

```
chap7data2 <- read.csv("chap7data2.csv", sep=";", header=T)
dim(chap7data2)
```

```
[1] 10 6
```

```
head(chap7data2)
```

	Detox	Enzyme1	Enzyme2	Enzyme3	Enzyme4	Enzyme5
1	56.250	106.329	90.756	94.650	162.791	114.737
2	75.000	144.726	203.361	131.687	255.814	112.632
3	115.625	136.287	672.269	123.457	191.860	153.684
4	68.750	154.430	183.193	113.169	133.721	116.842
5	96.875	385.232	140.336	117.284	174.419	87.368
6	168.750	593.544	146.218	152.263	273.256	94.737

```
chap7data3 <- read.csv("chap7data3.csv", sep=";", header=T)
dim(chap7data3)
```


```
[1] 13  5
```

```
head(chap7data3)
```


	X1	X2	X3	X4	Y
1	7	26	6	60	78.5
2	1	29	15	52	74.3
3	11	56	8	20	104.3
4	11	31	8	47	87.6
5	7	52	6	33	95.9
6	11	55	9	22	109.2

```
#scatter plot
```

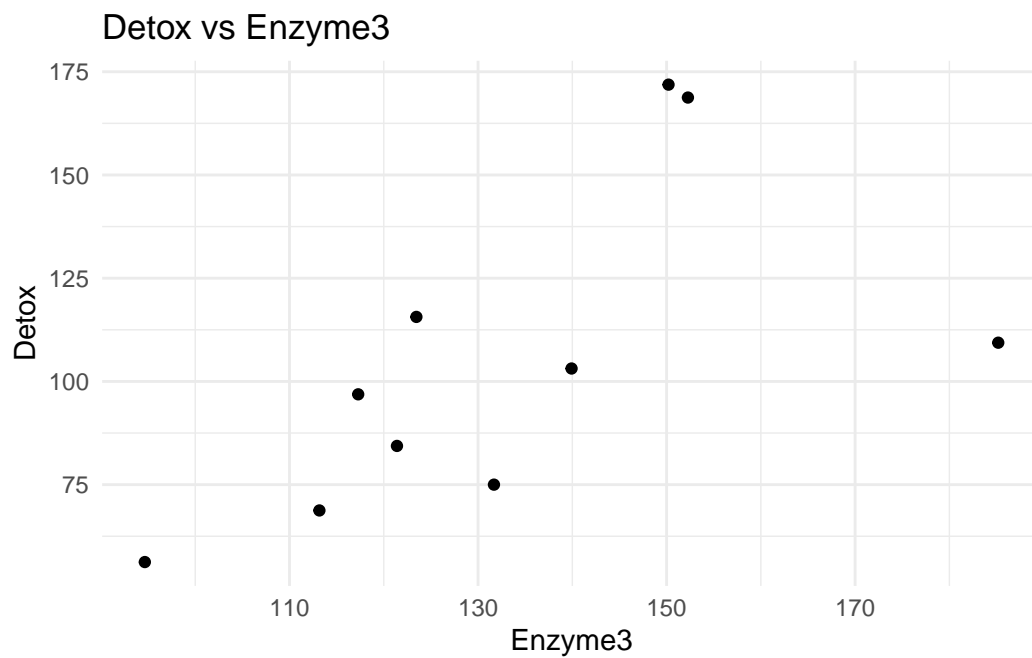
```
# Scatter plots of 'Detox' against each enzyme
ggplot(chap7data2, aes(x = Enzyme1, y = Detox)) + geom_point() + ggtitle("Detox vs Enzyme1")
  theme_minimal()
```




```
ggplot(chap7data2, aes(x = Enzyme2, y = Detox)) + geom_point() + ggtitle("Detox vs Enzyme2")  
theme_minimal()
```



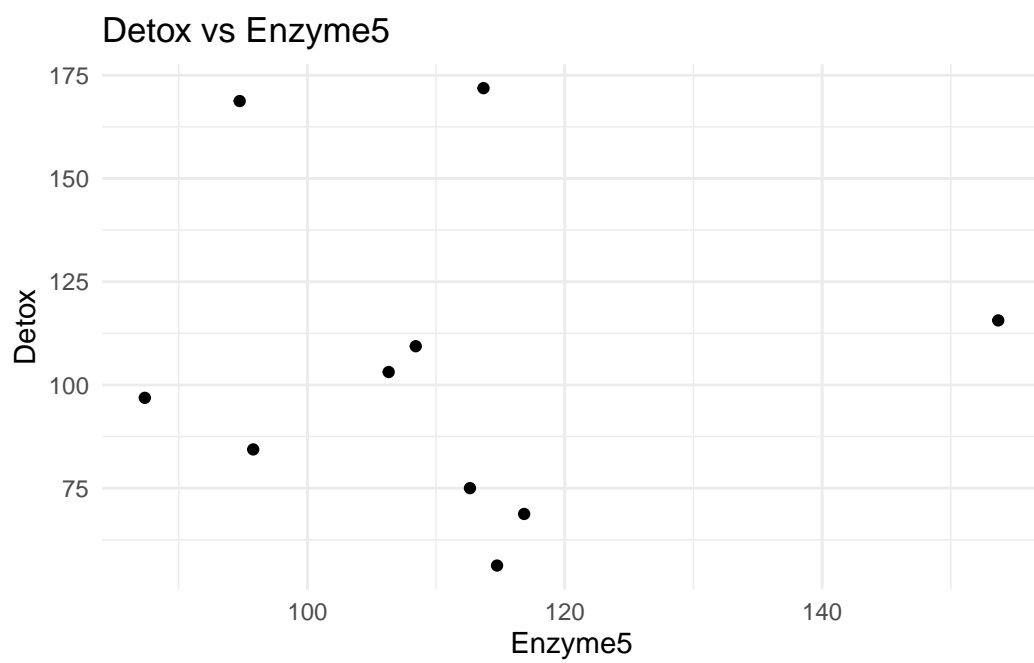
```
ggplot(chap7data2, aes(x = Enzyme3, y = Detox)) + geom_point() + ggtitle("Detox vs Enzyme3")
  theme_minimal()
```



```
ggplot(chap7data2, aes(x = Enzyme4, y = Detox)) + geom_point() + ggtitle("Detox vs Enzyme4")
  theme_minimal()
```




```
ggplot(chap7data2, aes(x = Enzyme5, y = Detox)) + geom_point() + ggtitle("Detox vs Enzyme5")
  theme_minimal()
```



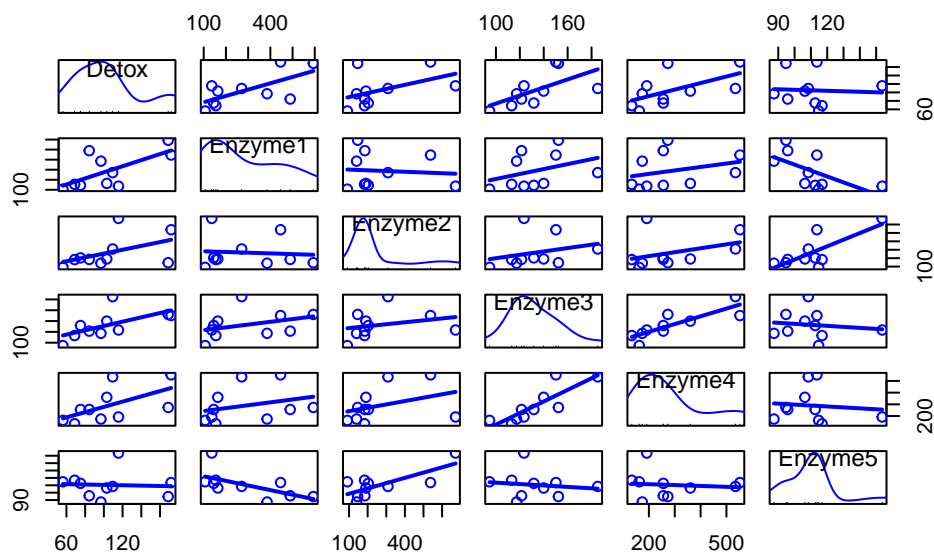
#or

```
plot(chap7data2)
```



#or

```
scatterplotMatrix(~Detox + Enzyme1 + Enzyme2 + Enzyme3 + Enzyme4 + Enzyme5,  
  reg.line = lm,  
  smooth = FALSE,  
  spread = FALSE,  
  span = 0.5,  
  ellipse = FALSE,  
  levels = c(0.5, 0.9),  
  id.n = 0,  
  diagonal = 'density',  
  data = chap7data2) |> suppressWarnings()
```



Interpretation

This document provides an interpretation of the scatterplot matrix that includes the response variable (**Detox**) and explanatory variables (**Enzyme1**, **Enzyme2**, **Enzyme3**, **Enzyme4**, and **Enzyme5**). The scatterplot matrix is used to visualize the relationships between these variables and identify potential correlations.

1. Relationships Between Detox and Each Enzyme

Examine row-wise.

- **Detox vs. Enzyme1:**

There is a positive linear relationship, as indicated by the upward trend in the scatter plot. The regression line is slightly upward, and the points are moderately scattered around it, suggesting a moderate positive correlation.

- **Detox vs. Enzyme2:**

The plot shows a positive linear relationship, as evidenced by the upward trend and a moderately steep regression line. The points are widely scattered without any discernible pattern, suggesting little to no linear relationship between **Detox** and **Enzyme2**.

- **Detox vs. Enzyme3:**

The plot shows a positive linear relationship, as evidenced by the upward trend and a moderately steep regression line. Points are somewhat dispersed, but there is a noticeable trend suggesting a moderate positive correlation.

- **Detox vs. Enzyme4:**

The plot shows a positive linear relationship, as evidenced by the upward trend and a moderately steep regression line. The points are widely scattered. This suggests a moderate correlation between Detox and Enzyme4.

- **Detox vs. Enzyme5:**

The scatter plot indicates a weak relationship with a relatively flat, slightly negative, regression line. The points are fairly scattered, with no clear trend, suggesting little to no linear correlation between Detox and Enzyme5.

2. Relationships Among Explanatory Variables (Enzyme1 to Enzyme5)

- **Enzyme1 vs. Enzyme2:**

There is a slight negative linear relationship with an downward trend. The regression line is mildly flat, and points are dispersed, indicating a slight negative correlation to little correlation.

- **Enzyme1 vs. Enzyme3, Enzyme4, Enzyme5:**

All these plots show relatively moderate positive linear relationships, with upward trends and points moderately clustered around the regression lines.

- **Enzyme1 vs. Enzyme5:**

There is a strong negative linear relationship with an downward trend. The regression line is steep and downwards sloping, and points are relatively clustered, indicating a strong negative correlation.

- **Enzyme2 vs. Enzyme3:**

The relationship appears to be non-linear, with points forming a curved pattern. This suggests a non-linear relationship between Enzyme2 and Enzyme3.

- **Enzyme2 vs. Enzyme4, Enzyme5:**

These plots show no clear pattern, suggesting little to no linear relationship between Enzyme2 and either Enzyme4 or Enzyme5.

- **Enzyme3 vs. Enzyme4:**

There seems to be a positive relationship with an upward trend in the scatter plot, though the points are somewhat scattered, indicating a moderate positive correlation.

- **Enzyme3 vs. Enzyme5:**

Shows a weak or no discernible relationship, with the points widely scattered around a flat regression line.

- **Enzyme4 vs. Enzyme5:**
No clear relationship is observed here, as the points are scattered without any noticeable pattern.

3. Diagonal Plots (Density Plots)

- The diagonal plots display the distribution of each variable. For example:
 - **Detox** appears to have a somewhat skewed distribution, possibly slightly left-skewed.
 - **Enzyme1** and **Enzyme3** appear to be normally distributed with slight variation.
 - **Enzyme2** shows a highly skewed distribution with a long tail.
 - **Enzyme4** and **Enzyme5** show more uniform distributions with less obvious skewness.

Conclusion

- **Detox** has a moderate positive relationship with **Enzyme1** and **Enzyme3**, suggesting these enzymes might be more influential in predicting **Detox** levels.
- **Enzyme2** and **Enzyme5** appear to have little to no linear relationship with **Detox**.
- The relationships among explanatory variables vary, with some moderate positive linear correlations (like **Enzyme1** with **Enzyme2**), and some non-linear or weak relationships.

Multiple linear regression

```
chap7data2.model11 <- lm(Detox~.,data=chap7data2)
# the `.` is shorthand to denote all predictor (i.e., X_j) variables
# to be included in the model

# Alternatively, you can specify the variables independently:
# chap7data2.model11 <- lm(Detox ~ Enzyme1 + Enzyme2 + Enzyme3 +
# Enzyme4 + Enzyme5, data=chap7data2)

summary(chap7data2.model11)
```

Call:

```
lm(formula = Detox ~ ., data = chap7data2)
```

Residuals:

1	2	3	4	5	6	7	8
2.7526	-3.0631	-0.8577	-2.1760	6.2696	21.6524	-36.4182	11.8682
9	10						
-19.7684	19.7406						

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-87.96535	189.66881	-0.464	0.667
Enzyme1	0.17063	0.10045	1.699	0.165
Enzyme2	0.03305	0.14039	0.235	0.825
Enzyme3	0.31844	0.67350	0.473	0.661
Enzyme4	0.04325	0.13468	0.321	0.764
Enzyme5	0.72460	1.69421	0.428	0.691

Residual standard error: 26.35 on 4 degrees of freedom
Multiple R-squared: 0.798, Adjusted R-squared: 0.5455
F-statistic: 3.161 on 5 and 4 DF, p-value: 0.1439

#the allSubsets() function

```
allSubsets <- function(data,y.name="Y",
                        perf.measure=c("adj.r.squared","r.squared",
                                       "AIC","perc.error"))
{
  Cols <- names(data)
  Cols <- Cols[! Cols %in% y.name]
  n <- length(Cols)

  id <- unlist(
    lapply(1:n,
           function(i)combn(1:n,i,simplify=F)
    )
    ,recursive=F)

  Formulas <- sapply(id,function(i)
    paste(y.name,"~",paste(Cols[i],collapse="+"))
  )

  result.mat <- matrix(0,nrow=length(Formulas),6)
```

```

result.mat[,1] <- Formulas

#get all adjusted R2 values
for(i in 1:length(Formulas)){
  result.mat[i,2] <- summary(
    lm(
      as.formula(Formulas[i]), data=data)
    )$adj.r.squared
}

#get all R2 values
for(i in 1:length(Formulas)){
  result.mat[i,3] <- summary(
    lm(
      as.formula(Formulas[i]), data=data)
    )$r.squared
}

#get all AIC's
for(i in 1:length(Formulas)){
  result.mat[i,4] <- AIC(
    lm(
      as.formula(Formulas[i]), data=data)
    )
}

#get all sigma's
for(i in 1:length(Formulas)){
  result.mat[i,5] <- summary(
    lm(
      as.formula(Formulas[i]), data=data
    )
  )$sigma/mean(data[,y.name])
}

#get all mallows cp's
for(i in 1:length(Formulas)){
  result.mat[i,6] <- mallows.cp(data=data,model=lm(
    as.formula(Formulas[i]), data=data), y.name=y.name)
}

colnames(result.mat) <- c("Model",

```

```

        "adj.r.squared",
        "r.squared",
        "AIC", "perc.error",
        "mallows.cp")

final.output <- data.frame(
  result.mat[order(result.mat[,perf.measure],decreasing=T),]
)

return(final.output)
}

```

calculating the Mallow's Cp statistic (function)

Goal: Choose models where $C_p \approx p$

Interpretation: Helps select models with a good balance of fit and complexity.

```

mallows.cp <- function(data,model,y.name)
{

  n <- nrow(data)
  p <- length(model$coefficients)

  mean.sq.error <- summary(
    lm(
      as.formula(paste(y.name,"~.")),data=data
    )
  )$sigma^2

  sse <- summary(model)$sigma^2*summary(model)$df[2]

  cp <- sse/mean.sq.error-(n-2*p)
  cp/100
}

```

Variable selection

```
stepwise(chap7data2.model1, direction='backward/forward', criterion='AIC')
```

Direction: backward/forward

Criterion: AIC

Start: AIC=68.27

Detox ~ Enzyme1 + Enzyme2 + Enzyme3 + Enzyme4 + Enzyme5

	Df	Sum of Sq	RSS	AIC
- Enzyme2	1	38.50	2816.7	66.407
- Enzyme4	1	71.60	2849.8	66.524
- Enzyme5	1	127.05	2905.2	66.717
- Enzyme3	1	155.27	2933.4	66.813
<none>			2778.2	68.270
- Enzyme1	1	2004.15	4782.3	71.701

Step: AIC=66.41

Detox ~ Enzyme1 + Enzyme3 + Enzyme4 + Enzyme5

	Df	Sum of Sq	RSS	AIC
- Enzyme3	1	130.4	2947.1	64.860
- Enzyme4	1	221.1	3037.8	65.163
<none>			2816.7	66.407
+ Enzyme2	1	38.5	2778.2	68.270
- Enzyme5	1	2216.4	5033.0	70.212
- Enzyme1	1	5478.4	8295.1	75.208

Step: AIC=64.86

Detox ~ Enzyme1 + Enzyme4 + Enzyme5

	Df	Sum of Sq	RSS	AIC
<none>			2947.1	64.860
+ Enzyme3	1	130.4	2816.7	66.407
+ Enzyme2	1	13.6	2933.4	66.813
- Enzyme4	1	1837.7	4784.8	67.706
- Enzyme5	1	2262.8	5209.8	68.557
- Enzyme1	1	6029.0	8976.1	73.997

Call:

```
lm(formula = Detox ~ Enzyme1 + Enzyme4 + Enzyme5, data = chap7data2)
```

Coefficients:

(Intercept)	Enzyme1	Enzyme4	Enzyme5
-102.4349	0.1938	0.1008	1.1069

#or

```
allSubsets(chap7data2,y.name="Detox",perf.measure="mallows.cp")
```

	Model	adj.r.squared
1	Detox ~ Enzyme5	-0.121740174648695
2	Detox ~ Enzyme3+Enzyme4+Enzyme5	0.0953356909980044
3	Detox ~ Enzyme2	0.121913336927653
4	Detox ~ Enzyme4+Enzyme5	0.16091426388108
5	Detox ~ Enzyme2+Enzyme3+Enzyme4	0.220935843219152
6	Detox ~ Enzyme3+Enzyme5	0.202661735249138
7	Detox ~ Enzyme3+Enzyme4	0.223753049711396
8	Detox ~ Enzyme2+Enzyme4	0.255553133684301
9	Detox ~ Enzyme4	0.265785404743971
10	Detox ~ Enzyme2+Enzyme3+Enzyme4+Enzyme5	0.374128029592512
11	Detox ~ Enzyme2+Enzyme3	0.329335344543953
12	Detox ~ Enzyme3	0.301375192960427
13	Detox ~ Enzyme1+Enzyme2+Enzyme3+Enzyme4+Enzyme5	0.54551810809289
14	Detox ~ Enzyme2+Enzyme4+Enzyme5	0.404467243817661
15	Detox ~ Enzyme1+Enzyme3+Enzyme4	0.451096799223505
16	Detox ~ Enzyme2+Enzyme3+Enzyme5	0.46270944285908
17	Detox ~ Enzyme1	0.384163301765189
18	Detox ~ Enzyme1+Enzyme2+Enzyme4+Enzyme5	0.616094255366656
19	Detox ~ Enzyme1+Enzyme2+Enzyme3+Enzyme4	0.619787416838814
20	Detox ~ Enzyme1+Enzyme2+Enzyme3+Enzyme5	0.627043441524617
21	Detox ~ Enzyme2+Enzyme5	0.476586035306095
22	Detox ~ Enzyme1+Enzyme3+Enzyme4+Enzyme5	0.631375639659066
23	Detox ~ Enzyme1+Enzyme3	0.505130867699246
24	Detox ~ Enzyme1+Enzyme4	0.512984256086067
25	Detox ~ Enzyme1+Enzyme2+Enzyme5	0.589864494082874
26	Detox ~ Enzyme1+Enzyme5	0.552720532095238
27	Detox ~ Enzyme1+Enzyme2+Enzyme4	0.658170511105981
28	Detox ~ Enzyme1+Enzyme3+Enzyme5	0.668698598042593
29	Detox ~ Enzyme1+Enzyme4+Enzyme5	0.678591090390398
30	Detox ~ Enzyme1+Enzyme2+Enzyme3	0.682256480210184

31		Detox ~ Enzyme1+Enzyme2 0.646870701721216		
	r.squared	AIC	perc.error	mallows.cp
1	0.0028976225344936	106.614682983029	0.39431982572297	0.137453882255527
2	0.396890460665336	105.587136855917	0.354116533733155	0.0994323899514434
3	0.219478521713469	104.165771062923	0.348876057988313	0.0945648667123251
4	0.347377760796396	104.376133095174	0.341040255742936	0.0892372756191775
5	0.480623895479435	104.092431403323	0.328616289784617	0.0828508511322706
6	0.379848016304885	103.865794310814	0.332448041968861	0.0828072658700512
7	0.396252371997752	103.597711242751	0.328021590110949	0.0795587491774218
8	0.420985770643346	103.179419011073	0.321232390178822	0.074660851334314
9	0.347364804216863	102.376331624279	0.319017114715357	0.0692398413807154
10	0.652293349773618	102.079740132558	0.294540640692888	0.0688555453531038
11	0.478371934645297	102.135696635339	0.304898428354942	0.0632968017381836
12	0.37900017152038	101.879456538239	0.311189172259715	0.0629751626156955
13	0.798008048041284	98.6484270290975	0.250992513755551	0.06
14	0.602978162545108	101.406061334488	0.287313017352848	0.0586213180485603
15	0.63406453281567	100.590718467776	0.27583562615466	0.0524653558987582
16	0.641806295239386	100.376887629136	0.272902227212976	0.0509322725558538
17	0.452589601569056	100.618136405545	0.292169756597092	0.0484024176453974
18	0.786719030759254	97.1922524896298	0.230682570166493	0.0422355380345724
19	0.788770787132675	97.0955870885837	0.22957030910281	0.0418292334558949
20	0.79280191195812	96.9029013083644	0.227369182566776	0.041030959111525
21	0.592900249682518	99.6567309027967	0.269355054023775	0.0406170238704733
22	0.795208688699481	96.7860631506573	0.226044785017069	0.0405543506688574
23	0.615101785988302	99.0959376823384	0.261907333353099	0.0362205048823659
24	0.62120997695583	98.9359687073529	0.259820836463091	0.0350109139242518
25	0.726576329388583	97.6763734731404	0.238432888587324	0.034145458363074
26	0.652115969407407	98.0848402690308	0.248995763195505	0.0288906715775875
27	0.772113674070654	95.8546178223231	0.21767434129564	0.0251278031069124
28	0.779132398695062	95.5417828096502	0.214296023808806	0.0237379012704586
29	0.785727393593599	95.2386391415135	0.211072391978277	0.0224319096535481
30	0.788170986806789	95.1239425579172	0.209865392142655	0.0219480105299455
31	0.725343879116501	95.7213469375342	0.221242904754148	0.0143895176456604

Interpretation 1

The stepwise regression analysis identified that the best model for predicting **Detox** includes the variables **Enzyme1**, **Enzyme4**, and **Enzyme5**. This model was chosen because it has the lowest Akaike Information Criterion (AIC) value (64.86), indicating the best balance between model complexity and fit.

- **Enzyme1** and **Enzyme4** have small positive effects on **Detox**, while **Enzyme5** has a

larger positive effect.

- The variables `Enzyme2` and `Enzyme3` were excluded from the final model as they did not significantly improve the model's fit.

This suggests that `Enzyme1`, `Enzyme4`, and `Enzyme5` are the most important predictors of `Detox` in this dataset.

Interpretation 2

Based on the regression analysis and the adjusted R-squared values for various models, the following conclusions can be drawn:

1. **Best Model:** The model `Detox ~ Enzyme1 + Enzyme2 + Enzyme3` has the highest adjusted R-squared value of 0.6823. This indicates that it explains about 68.23% of the variability in `Detox` and is the best model among all those considered in terms of fit.
2. **Other Notable Models:**
 - The models `Detox ~ Enzyme1 + Enzyme2 + Enzyme5` (Adjusted R-squared: 0.6687) and `Detox ~ Enzyme1 + Enzyme2 + Enzyme4` (Adjusted R-squared: 0.6582) also have high adjusted R-squared values. These models provide a good fit and are close in performance to the best model.
3. **Poorly Performing Models:**
 - Models with adjusted R-squared values close to or below 0, such as `Detox ~ Enzyme5` (-0.1217) and `Detox ~ Enzyme3` (0.3014), indicate a poor fit. These models do not explain much of the variability in `Detox` and are not suitable for prediction.
4. **Model Selection:**
 - The top models to consider for predicting `Detox` are:
 - `Detox ~ Enzyme1 + Enzyme2 + Enzyme3` (Adjusted R-squared: 0.6823)
 - `Detox ~ Enzyme1 + Enzyme2 + Enzyme5` (Adjusted R-squared: 0.6687)
 - `Detox ~ Enzyme1 + Enzyme2 + Enzyme4` (Adjusted R-squared: 0.6582)
 - These models provide a good balance between complexity (number of variables) and explanatory power.

Conclusion:

The best model based on adjusted R-squared is `Detox ~ Enzyme1 + Enzyme2 + Enzyme3`, which explains 68.23% of the variance in `Detox`. Models with fewer variables generally have lower adjusted R-squared values and may not provide sufficient explanatory power, while models with negative adjusted R-squared values should be discarded as they do not fit the data well.

Another example with data 3

```
chap7data3.model11 <- lm(Y~.,data=chap7data3)
#chap7data3.model11 <- lm(Y~X1+X2+X3+X4,data=chap7data3)
summary(chap7data3.model11)
```

Call:

```
lm(formula = Y ~ ., data = chap7data3)
```

Residuals:

Min	1Q	Median	3Q	Max
-3.1750	-1.6709	0.2508	1.3783	3.9254

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	62.4054	70.0710	0.891	0.3991
X1	1.5511	0.7448	2.083	0.0708
X2	0.5102	0.7238	0.705	0.5009
X3	0.1019	0.7547	0.135	0.8959
X4	-0.1441	0.7091	-0.203	0.8441

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.446 on 8 degrees of freedom

Multiple R-squared: 0.9824, Adjusted R-squared: 0.9736


F-statistic: 111.5 on 4 and 8 DF, p-value: 4.756e-07

```
# Test for multicollinearity and variable influence using Variable Influence Factor [vif()]
# VIF equal to 1 = variables are not correlated
# VIF between 1 and 5 = variables are moderately correlated
```

```
# VIF greater than 5 = variables are highly correlated
vif(chap7data3.model1)
```

```
      X1      X2      X3      X4
38.49621 254.42317 46.86839 282.51286
```

```
base::plot(chap7data3)
```



```
stepwise(chap7data3.model1, direction='backward', criterion='AIC')
```

```
Direction: backward
Criterion: AIC
```

```
Start: AIC=26.94
Y ~ X1 + X2 + X3 + X4
```

	Df	Sum of Sq	RSS	AIC
- X3	1	0.1091	47.973	24.974
- X4	1	0.2470	48.111	25.011
- X2	1	2.9725	50.836	25.728

```

<none>                47.864 26.944
- X1      1    25.9509 73.815 30.576

```

```

Step:  AIC=24.97
Y ~ X1 + X2 + X4

```

```

      Df Sum of Sq    RSS    AIC
<none>                47.97 24.974
- X4      1      9.93  57.90 25.420
- X2      1     26.79  74.76 28.742
- X1      1    820.91 868.88 60.629

```

```

Call:
lm(formula = Y ~ X1 + X2 + X4, data = chap7data3)

```

```

Coefficients:
(Intercept)          X1          X2          X4
    71.6483      1.4519      0.4161     -0.2365

```

```

stepwise(chap7data3.model1, direction='forward', criterion='AIC')

```

```

Direction:  forward
Criterion:  AIC

```

```

Start:  AIC=71.44
Y ~ 1

```

```

      Df Sum of Sq    RSS    AIC
+ X4      1    1831.90  883.87 58.852
+ X2      1    1809.43  906.34 59.178
+ X1      1    1450.08 1265.69 63.519
+ X3      1     776.36 1939.40 69.067
<none>                2715.76 71.444

```

```

Step:  AIC=58.85
Y ~ X4

```

```

      Df Sum of Sq    RSS    AIC
+ X1      1     809.10  74.76 28.742
+ X3      1     708.13 175.74 39.853

```

```

<none>                883.87 58.852
+ X2      1      14.99 868.88 60.629

```

```

Step:  AIC=28.74
Y ~ X4 + X1

```

```

      Df Sum of Sq    RSS    AIC
+ X2    1    26.789 47.973 24.974
+ X3    1    23.926 50.836 25.728
<none>                74.762 28.742

```

```

Step:  AIC=24.97
Y ~ X4 + X1 + X2

```

```

      Df Sum of Sq    RSS    AIC
<none>                47.973 24.974
+ X3    1    0.10909 47.864 26.944

```

```

Call:
lm(formula = Y ~ X4 + X1 + X2, data = chap7data3)

```

Coefficients:

```

(Intercept)      X4      X1      X2
  71.6483    -0.2365    1.4519    0.4161

```

```

stepwise(chap7data3.model1, direction='backward/forward', criterion='AIC')

```

```

Direction:  backward/forward
Criterion:  AIC

```

```

Start:  AIC=26.94
Y ~ X1 + X2 + X3 + X4

```

```

      Df Sum of Sq    RSS    AIC
- X3    1    0.1091 47.973 24.974
- X4    1    0.2470 48.111 25.011
- X2    1    2.9725 50.836 25.728
<none>                47.864 26.944
- X1    1   25.9509 73.815 30.576

```

Step: AIC=24.97
Y ~ X1 + X2 + X4

	Df	Sum of Sq	RSS	AIC
<none>			47.97	24.974
- X4	1	9.93	57.90	25.420
+ X3	1	0.11	47.86	26.944
- X2	1	26.79	74.76	28.742
- X1	1	820.91	868.88	60.629

Call:

```
lm(formula = Y ~ X1 + X2 + X4, data = chap7data3)
```

Coefficients:

(Intercept)	X1	X2	X4
71.6483	1.4519	0.4161	-0.2365

```
allSubsets(chap7data3,y.name="Y",perf.measure="adj.r.squared")
```

	Model	adj.r.squared	r.squared	AIC
1	Y ~ X1+X2+X4	0.976447268267236	0.982335451200427	63.8662854718626
2	Y ~ X1+X2+X3	0.976379572292115	0.982284679219087	63.9035968705951
3	Y ~ X1+X3+X4	0.975041456783125	0.981281092587343	64.6199522325816
4	Y ~ X1+X2	0.974414049442758	0.978678374535632	64.3123927621906
5	Y ~ X1+X2+X3+X4	0.97356343061152	0.98237562040768	65.8366897916517
6	Y ~ X1+X4	0.966965257260317	0.972471047716931	67.6341062672404
7	Y ~ X2+X3+X4	0.963759945848363	0.972819959386272	69.468286610933
8	Y ~ X3+X4	0.922347568738969	0.935289640615807	78.744985808948
9	Y ~ X2+X3	0.816430499370665	0.847025416142221	89.9295421295082
10	Y ~ X4	0.644954869961756	0.674541964131609	97.7440447788562
11	Y ~ X2	0.63592900832723	0.666268257633294	98.0703964203777
12	Y ~ X2+X4	0.61607248954006	0.68006040795005	99.5217275129777
13	Y ~ X1	0.491579662365491	0.533948023835033	102.41187201392
14	Y ~ X1+X3	0.457800098613821	0.548166748844851	104.009080451406
15	Y ~ X3	0.220952070432522	0.285872731229812	107.959804394717
	perc.error	mallows.cp		
1	0.0241948282253226	0.0301823347348734		
2	0.0242295740305319	0.0304127972306419		
3	0.0249064320090681	0.0349682444234851		
4	0.0252175376870345	0.0267824159831842		
5	0.0256332957861164	0.05		

```

6 0.0286541391065289 0.0549585082475868
7 0.0300120954747619 0.0733747399565644
8 0.0439318436492178 0.223731119646975
9 0.0675463931752181 0.624377163435314
10 0.0939385128212314 1.38730833491677
11 0.095125058915528 1.42486406936961
12 0.0976846931820687 1.38225919754642
13 0.112412176793213 2.0254876912345
14 0.116086473289248 1.98094652569577
15 0.139150249457931 3.15154284140081

```

```

chap7data3.model2 <- lm(Y ~ X1 + X2 + X4,data=chap7data3)
summary(chap7data3.model2)

```

Call:

```
lm(formula = Y ~ X1 + X2 + X4, data = chap7data3)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-3.0919	-1.8016	0.2562	1.2818	3.8982

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	71.6483	14.1424	5.066	0.000675 ***
X1	1.4519	0.1170	12.410	5.78e-07 ***
X2	0.4161	0.1856	2.242	0.051687 .
X4	-0.2365	0.1733	-1.365	0.205395

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.309 on 9 degrees of freedom

Multiple R-squared: 0.9823, Adjusted R-squared: 0.9764

F-statistic: 166.8 on 3 and 9 DF, p-value: 3.323e-08

```
vif(chap7data3.model2)
```

	X1	X2	X4
	1.06633	18.78031	18.94008

```
chap7data3.model3 <- lm(Y ~ X1 + X2,data=chap7data3)
summary(chap7data3.model3)
```

Call:

```
lm(formula = Y ~ X1 + X2, data = chap7data3)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.893	-1.574	-1.302	1.363	4.048

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	52.57735	2.28617	23.00	5.46e-10 ***
X1	1.46831	0.12130	12.11	2.69e-07 ***
X2	0.66225	0.04585	14.44	5.03e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.406 on 10 degrees of freedom

Multiple R-squared: 0.9787, Adjusted R-squared: 0.9744

F-statistic: 229.5 on 2 and 10 DF, p-value: 4.407e-09

```
vif(chap7data3.model3)
```

X1	X2
1.055129	1.055129

#Assess normality

```
shapiro.test(chap7data3.model3$residuals)
```

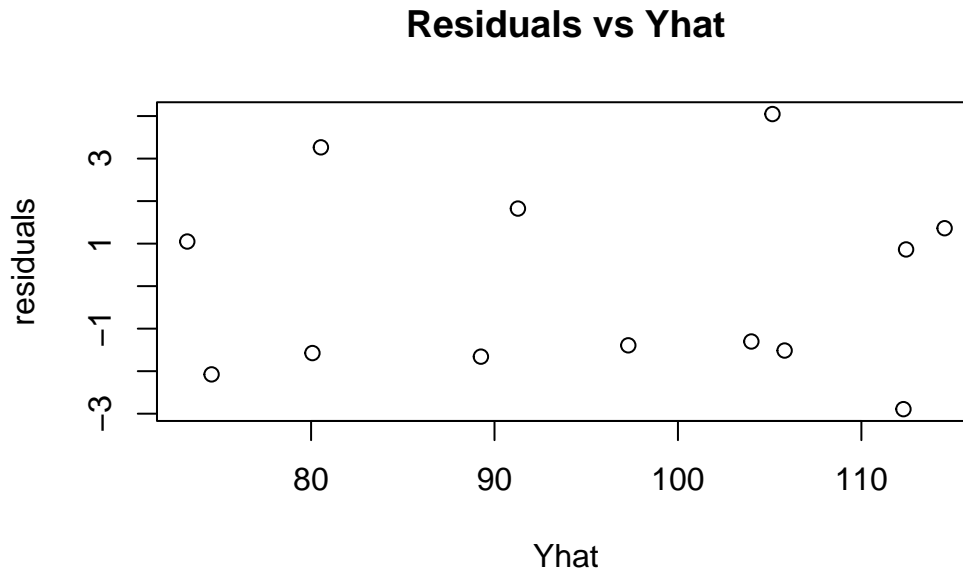
Shapiro-Wilk normality test

data: chap7data3.model3\$residuals

W = 0.90527, p-value = 0.158

#Assess homoscedasticity


```
plot(y=chap7data3.model3$residuals,x=chap7data3.model3$fitted.values,
     main="Residuals vs Yhat",ylab="residuals",xlab="Yhat")
```



Exercises

Exercise 1

Use the `swiss` dataset in R (do this by running the code `data(swiss)`) to fit a multiple linear regression model of **Fertility** on the other variables, and find the best fitting model according to Mallows' C_p criterion. Do this using all department observations (i.e., rownames) except Moutier (i.e. the 4th observation/row).

To know more about the dataset, you can type and run `?swiss` in the R console.

For this final model:

- Test the model assumptions
- Interpret the regression coefficients (parameter estimates)
- Interpret the coefficient of determination (R^2)

- Predict Fertility for Moutier using the collected data for the department. What is the prediction? How does the predicted value and observed value compare to the predicted interval?

Exercise 2 and further reading

Polynomial regression

See pdf file Biometry 721-821 Workbook (2024), Chapter 7, Section 2, pages 52-54 on polynomial regression.

With the data from C&S Example 8.1, p. 90 (chap8exer1.csv): Fit a quadratic regression model of wheat yield (Y) on nitrogen fertilizer (X) and give the equation of the regression line, as well as the coefficient of determination R^2 . Also check that the assumptions of this polynomial regression model are justified. Does the quadratic model fit the data better than the simple linear regression (straight-line) model?